Global Workspace Theory (GWT)

A Comprehensive Mathematical & Conceptual Guide

From the Theater Metaphor to Computational Models
of Conscious Access and Broadcast

Baars • Dehaene • Changeux • Mashour

Covering Prerequisites • Core Principles • Neural Mechanisms • Formal Models

Comprehensive Study Material | February 2026

Contents

1 Introduction to Global Workspace Theory **3**
1.1 What is Global Workspace Theory? . . . 3
1.2 Historical Context and Motivation . . . 3
1.3 The Hard Problem and GWT's Stance . . . 4
1.4 Versions and Evolution of the Theory . . . 4

2 Mathematical and Conceptual Prerequisites **5**
2.1 Dynamical Systems Theory . . . 5
2.2 Neural Population Models . . . 5
2.2.1 Wilson–Cowan Model . . . 5
2.2.2 Mean-Field Approximations . . . 6
2.3 Information Theory for Neural Coding . . . 6
2.3.1 Mutual Information in Neural Coding . . . 6
2.3.2 Transfer Entropy . . . 7
2.3.3 Granger Causality . . . 7
2.4 Signal Detection Theory . . . 7
2.5 Graph Theory for Brain Networks . . . 7

3 Core Principles of GWT **9**
3.1 Principle 1: Modularity and Specialization . . . 9
3.2 Principle 2: The Global Workspace as Integrative Hub . . . 9
3.3 Principle 3: Competition and Selection . . . 10
3.4 Principle 4: Ignition and Broadcast . . . 10
3.5 Principle 6: Conscious vs. Unconscious Processing . . . 11

4 The Global Neuronal Workspace (GNWT) **12**
4.1 Architecture of the Global Neuronal Workspace . . . 12
4.1.1 Workspace Neurons . . . 12
4.1.2 Peripheral Processors . . . 12
4.1.3 Connectivity Architecture . . . 12
4.2 The Ignition Dynamics: Mathematical Model . . . 13
4.2.1 Two-Population Model . . . 13
4.2.2 Ignition as a Bifurcation . . . 14
4.2.3 Top-Down Modulation and Attention . . . 14
4.3 Multi-Area Spiking Model . . . 14

5 Formal Computational Models **16**
5.1 The Broadcast Formalism . . . 16
5.1.1 Competition via Gating . . . 16
5.2 Information-Theoretic Characterization . . . 16
5.2.1 Broadcast as Information Gain . . . 16
5.2.2 Workspace Capacity and Information Bottleneck . . . 17
5.3 Bayesian Formulation . . . 17

6 Neural Signatures of Conscious Access **18**
6.1 The P3b / P300 Wave . . . 18
6.2 Late Sustained Activity . . . 18
6.3 Gamma-Band Synchrony and Long-Range Connectivity . . . 18
6.4 The "All-or-None" Pattern . . . 19

7 Key Experimental Paradigms **20**
7.1 Visual Masking . . . 20
7.2 The Attentional Blink . . . 20
7.3 Binocular Rivalry . . . 20
7.4 Inattentional Blindness and Change Blindness . . . 21

8 Consciousness Under Altered States **22**
8.1 General Anesthesia . . . 22
8.2 Sleep and Dreaming . . . 22
8.3 Disorders of Consciousness . . . 22

9 GWT vs. Other Theories **23**
9.1 GWT vs. Integrated Information Theory (IIT) . . . 23
9.2 GWT vs. Higher-Order Theories . . . 23
9.3 GWT vs. Recurrent Processing Theory . . . 24
9.4 GWT vs. Predictive Processing . . . 24

10 Criticisms and Open Questions **25**
10.1 Common Critiques . . . 25
10.2 Open Questions . . . 25

11 Further Reading and References **27**

Chapter 1

Introduction to Global Workspace Theory

1.1 What is Global Workspace Theory?

Global Workspace Theory (GWT) is one of the most influential scientific theories of consciousness, originally proposed by cognitive psychologist **Bernard Baars** in 1988. GWT provides a cognitive architecture for understanding how information becomes conscious and how conscious contents are distributed across the brain for widespread processing.

The central metaphor is that of a **theater**: consciousness is like a bright spotlight on a stage, illuminating one act at a time, while a vast audience of specialized unconscious processors sits in the darkened auditorium. Only information that enters the "spotlight"—the **global workspace**—becomes conscious, and once there, it is *broadcast* to all the audience members simultaneously.

Core Claim of GWT

Consciousness arises when information gains access to a **global workspace**—a shared, capacity-limited cognitive resource that enables the widespread **broadcasting** of information to multiple specialized, otherwise independent brain modules. Conscious access $\equiv$ global broadcast.

1.2 Historical Context and Motivation

GWT emerged from several converging intellectual traditions:

(a) Cognitive science and modularity: Jerry Fodor's (1983) influential proposal that the mind consists of specialized, informationally encapsulated **modules** (for language, face recognition, etc.) raised the question: how do these modules communicate and coordinate? GWT provides the answer: through a shared global workspace.

(b) Artificial intelligence: The "blackboard architecture" in early AI systems (Erman et al., 1980) provided a direct computational inspiration. In blackboard systems, multiple specialist modules read from and write to a shared working memory (the "blackboard"), enabling cooperative problem-solving. GWT adapts this idea to the brain.

(c) Attention and working memory: Decades of research on selective attention (Broadbent, Treisman) and working memory (Baddeley) showed that consciousness is associated with a capacity-limited bottleneck. GWT identifies this bottleneck with the global workspace.

(d) Neuroscience: Advances in neuroimaging and electrophysiology in the 1990s–2000s revealed that conscious perception is associated with widespread, sustained brain activation, while unconscious processing remains localized—exactly as GWT predicts.

1.3 The Hard Problem and GWT's Stance

Unlike IIT, GWT does not directly attempt to solve the Hard Problem of consciousness (why there is subjective experience at all). Instead, GWT focuses on the "**easy problems**" (in Chalmers' terminology)—specifically, the problem of **conscious access**: what determines which information becomes conscious and which remains unconscious?

GWT is thus primarily a theory of the **function** of consciousness: what consciousness does, how it works, and what neural mechanisms underlie it. It explains the *correlates* and *conditions* of consciousness without claiming to explain the intrinsic nature of subjective experience.

GWT vs. IIT on the Hard Problem

IIT claims to address the Hard Problem by identifying consciousness *with* integrated information. GWT is more modest: it characterizes the functional and neural architecture that determines conscious access, leaving the question of *why* this architecture gives rise to experience as an open (and perhaps separate) issue. Some theorists view these as complementary rather than competing.

1.4 Versions and Evolution of the Theory

GWT has evolved significantly since Baars' original proposal:

1. **GWT (Baars, 1988):** The original cognitive-level theory, formulated using the theater metaphor and concepts from cognitive science.
2. **Global Neuronal Workspace Theory (GNWT)** (Dehaene, Kerszberg & Changeux, 1998; Dehaene & Changeux, 2011): A neurally explicit version that maps GWT onto specific brain circuits, especially prefrontal-parietal networks connected by long-range excitatory neurons.
3. **Cognitive-Dynamic GWT (Mashour et al., 2020):** An updated framework incorporating temporal dynamics, multistability, and the role of anesthesia.
4. **Computational models:** Formal mathematical and simulation-based implementations, including spiking neural network models by Dehaene and colleagues.

Chapter 2

Mathematical and Conceptual Prerequisites

This chapter reviews the mathematical and conceptual foundations needed to understand the formal models of GWT.

2.1 Dynamical Systems Theory

GWT's computational models are built on dynamical systems. A **dynamical system** describes how a state evolves over time according to a rule.

Continuous Dynamical System

A continuous dynamical system is defined by a set of ordinary differential equations (ODEs):

$$\frac{d\mathbf{x}}{dt} = f(\mathbf{x}, t)$$

where $\mathbf{x}(t) \in \mathbb{R}^n$ is the state vector (e.g., firing rates of neural populations) and f specifies the dynamics. In neural models, f encodes the connectivity, excitation/inhibition balance, and nonlinear activation of neural populations.

Key concepts relevant to GWT:

- **Attractors:** Stable states toward which the system evolves. In GWT, a conscious percept may correspond to the system settling into a global attractor.
- **Bifurcations:** Qualitative changes in system behavior as parameters change. The transition from unconscious to conscious processing may involve a bifurcation (sudden ignition).
- **Bistability / Multistability:** The coexistence of multiple stable states. The global workspace may switch between competing representations.

2.2 Neural Population Models

The neural implementations of GWT use population-level models, where the fundamental variable is the average firing rate or membrane potential of a neural *population* (a group of similar neurons).

2.2.1 Wilson–Cowan Model

The Wilson–Cowan equations describe the interaction between excitatory (E) and inhibitory (I) neural populations:

$$\tau_E \frac{dE}{dt} = -E + S_E\big(w_{EE}E - w_{EI}I + I_{\text{ext}}^E\big) \tag{2.1}$$

$$\tau_I \frac{dI}{dt} = -I + S_I\big(w_{IE}E - w_{II}I + I_{\text{ext}}^I\big) \tag{2.2}$$

where τ_E, τ_I are time constants, w_{XY} are connection weights, I_{ext} is external input, and $S(\cdot)$ is a sigmoidal activation function:

$$S(x) = \frac{1}{1 + e^{-\beta(x-\theta)}} - \frac{1}{1 + e^{\beta\theta}}$$

with gain β and threshold θ. These equations form the building blocks for multi-area models of the global workspace.

2.2.2 Mean-Field Approximations

For large networks, **mean-field theory** replaces the detailed behavior of individual neurons with population-averaged quantities. For a population of N neurons with mean firing rate ν:

$$\nu = \Phi\big(\mu_V,\ \sigma_V^2\big)$$

where μ_V and σ_V^2 are the mean and variance of the membrane potential, and Φ is a transfer function derived from single-neuron properties (e.g., the frequency-current or f-I curve). The inputs to each population include:

$$\mu_V = \sum_j w_j \nu_j + I_{\text{ext}}, \qquad \sigma_V^2 = \sum_j w_j^2 \nu_j \tau_{\text{syn}}$$

where the sum is over all presynaptic populations j.

2.3 Information Theory for Neural Coding

Information theory provides tools for quantifying how much information neural activity carries about stimuli.

2.3.1 Mutual Information in Neural Coding

The mutual information between a stimulus S and a neural response R:

$$I(S;R) = H(R) - H(R \mid S) = \sum_{s,r} p(s,r) \log_2 \frac{p(s,r)}{p(s)\,p(r)}$$

In the context of GWT, this quantifies how much information about a stimulus is carried by neural activity in a given brain area. Conscious access (broadcast) should increase $I(S;R)$ for *distant* brain areas.

2.3.2 Transfer Entropy

Transfer entropy quantifies directed information flow between time series—essential for measuring broadcast dynamics:

$$T_{X\to Y} = \sum p\left(y_{t+1}, y_t^{(k)}, x_t^{(l)}\right) \log_2 \frac{p\left(y_{t+1} \mid y_t^{(k)}, x_t^{(l)}\right)}{p\left(y_{t+1} \mid y_t^{(k)}\right)}$$

where $y_t^{(k)}$ and $x_t^{(l)}$ denote the past k and l values of Y and X respectively. High transfer entropy from workspace regions to sensory areas indicates global broadcast.

2.3.3 Granger Causality

A linear counterpart to transfer entropy, **Granger causality** tests whether past values of X improve the prediction of Y beyond Y's own past:

$$Y_t = \sum_{i=1}^{p} a_i Y_{t-i} + \sum_{j=1}^{q} b_j X_{t-j} + \varepsilon_t$$

X Granger-causes Y if the coefficients b_j are jointly significant (F-test). This is widely used in neuroimaging to measure effective connectivity and information flow during conscious vs. unconscious processing.

2.4 Signal Detection Theory

GWT's empirical predictions are often tested using paradigms from **signal detection theory** (SDT), which provides a mathematical framework for analyzing perceptual decisions.

Signal Detection: d' and Criterion c

Given a signal+noise distribution $\mathcal{N}(\mu_S, \sigma^2)$ and a noise-only distribution $\mathcal{N}(\mu_N, \sigma^2)$:

$$d' = \frac{\mu_S - \mu_N}{\sigma}$$

measures discriminability (sensitivity), independent of response bias. The criterion:

$$c = -\frac{1}{2}\left[z(\text{hit rate}) + z(\text{false alarm rate})\right]$$

measures the observer's decision threshold. In GWT, conscious access corresponds to stimuli exceeding a *threshold* for global ignition, and SDT provides tools for separating sensitivity from criterion effects.

2.5 Graph Theory for Brain Networks

The global workspace is realized by a network of brain regions. Graph-theoretic measures characterize its structure.

Key Network Measures

For a graph $G = (V, E)$ with adjacency matrix A:

- **Degree:** $k_i = \sum_j A_{ij}$ — number of connections of node i.
- **Clustering coefficient:** $C_i = \frac{2|\{e_{jk}:v_j,v_k \in N_i, e_{jk} \in E\}|}{k_i(k_i-1)}$ — local connectivity.
- **Path length:** $L = \frac{1}{n(n-1)} \sum_{i \neq j} d(i, j)$ — average shortest path.
- **Small-world index:** $\sigma = \frac{C/C_{\text{random}}}{L/L_{\text{random}}}$ — high clustering with short paths.
- **Hub nodes:** Nodes with disproportionately high degree or betweenness centrality.

The global workspace is hypothesized to rely on **hub nodes** with high betweenness centrality that enable efficient broadcast.

Chapter 3

Core Principles of GWT

This chapter details the foundational principles of GWT. Unlike IIT's axiom-based structure, GWT is organized around a set of functional principles derived from cognitive science and neuroscience.

3.1 Principle 1: Modularity and Specialization

Principle 1 — Modularity

The brain consists of a large number of specialized, semi-autonomous **processors** (modules), each handling a specific type of information (e.g., color, motion, faces, language syntax). These modules operate largely in **parallel** and **unconsciously**.

Examples of specialized modules include the fusiform face area (FFA) for face processing, area V4 for color, area MT/V5 for motion, Wernicke's area for language comprehension, and the amygdala for threat detection. Each module can process information rapidly and automatically without requiring conscious awareness.

The key insight: most of what the brain does is *unconscious*. The vast majority of neural processing occurs in specialized modules that never enter awareness. Consciousness is the exception, not the rule.

3.2 Principle 2: The Global Workspace as Integrative Hub

Principle 2 — The Global Workspace

There exists a **global workspace**—a capacity-limited, shared cognitive resource—that allows information from any one module to be made available to **all other modules simultaneously**. This workspace acts as a central integration and broadcasting hub.

The global workspace is not a single brain region but a **distributed network** of neurons, primarily involving long-range cortical connections. In the neural implementation (GNWT), it is associated with:

- Prefrontal cortex (especially dorsolateral PFC)
- Parietal cortex (especially posterior parietal regions)
- Anterior cingulate cortex (ACC)
- Long-range cortico-cortical axons connecting these areas

The workspace is **capacity-limited**: only one coherent representation (or a very small number) can occupy the workspace at a time. This explains the well-known bottleneck of consciousness—

we can only be conscious of one thing (or a very few things) at a time.

3.3 Principle 3: Competition and Selection

Principle 3 — Competition for Access

Multiple unconscious processors compete for access to the global workspace. This competition is governed by **bottom-up salience** (stimulus strength, novelty, emotional valence) and **top-down attention** (goals, expectations, task relevance). Only the "winning" representation gains access.

The competition can be modeled as a winner-take-all (WTA) dynamics:

$$\tau \frac{dx_i}{dt} = -x_i + f\left(\sum_j w_{ij}^{\text{local}} x_j + I_i^{\text{ext}} - \alpha \sum_{k \neq i} x_k\right)$$

where x_i is the activity of the i-th competing representation, w^{local} are local recurrent connections (self-amplification), I^{ext} is external drive, α controls lateral inhibition (competition), and f is a nonlinear activation function. The representation with the strongest combined bottom-up and top-down support wins the competition and enters the workspace.

3.4 Principle 4: Ignition and Broadcast

Principle 4 — Ignition

When a representation wins the competition and enters the global workspace, it triggers a sudden, self-reinforcing, nonlinear **ignition**—a rapid amplification of activity through recurrent loops between workspace neurons. This ignition constitutes the **neural correlate of conscious access**.

Ignition is the signature event in GNWT. It is characterized by:

- A sudden, **nonlinear** transition from localized to widespread activity
- Self-sustaining activity through positive feedback loops
- **All-or-none** character: ignition either occurs fully or not at all
- A time delay (∼200–300 ms after stimulus onset)
- Associated electrophysiological signatures (P3b wave, late sustained activity, gamma synchrony)

Principle 5 — Global Broadcast

Following ignition, the contents of the global workspace are **broadcast** to all connected modules simultaneously. This broadcast makes the information available for diverse cognitive operations: verbal report, memory encoding, planning, emotional evaluation, and voluntary action.

The broadcast is the functional essence of consciousness in GWT: it is the mechanism by which a

piece of information transitions from being processed locally and unconsciously to being *globally available* and conscious.

3.5 Principle 6: Conscious vs. Unconscious Processing

Principle 6 — The Conscious/Unconscious Divide

Unconscious processing is: parallel, fast, automatic, capacity-unlimited, localized, and modular. **Conscious** processing is: serial (one content at a time), relatively slow, effortful, capacity-limited, distributed, and integrated.

This leads to a clear taxonomy:

Property	**Unconscious Processing**	**Conscious Processing**
Capacity	Unlimited (many modules)	Severely limited
Speed	Fast (<200 ms)	Slow (>300 ms)
Processing mode	Parallel	Serial
Neural scope	Local, modular	Global, distributed
Flexibility	Fixed, automatic	Flexible, voluntary
Reportability	Not reportable	Reportable
Information sharing	Limited to local circuits	Broadcast to all modules

Chapter 4

The Global Neuronal Workspace (GNWT)

This chapter covers the neurally explicit version of GWT, developed primarily by Stanislas Dehaene and Jean-Pierre Changeux.

4.1 Architecture of the Global Neuronal Workspace

4.1.1 Workspace Neurons

The global neuronal workspace consists of a specific population of neurons characterized by:

Workspace Neuron Properties

1. **Long-range axons:** Workspace neurons have long-range cortico-cortical axons that project across distant brain areas. They are concentrated in layers II/III of the cortex.
2. **Prefrontal-parietal distribution:** They are especially dense in prefrontal and parietal cortex but are found throughout the cortex.
3. **Recurrent connectivity:** Workspace neurons form dense recurrent loops with each other, enabling self-sustaining activity.
4. **Top-down projections:** They project back to sensory areas, modulating local processing through feedback connections.

4.1.2 Peripheral Processors

Contrasted with workspace neurons are the **peripheral processors**—specialized cortical and subcortical areas that perform specific computations (visual processing, auditory analysis, motor planning, etc.). These processors:

- Have primarily local, short-range connections
- Operate automatically and in parallel
- Can process information without workspace involvement (unconscious processing)
- Compete for access to the workspace via bottom-up signals

4.1.3 Connectivity Architecture

The GNWT architecture can be summarized as a **hub-and-spoke** model:

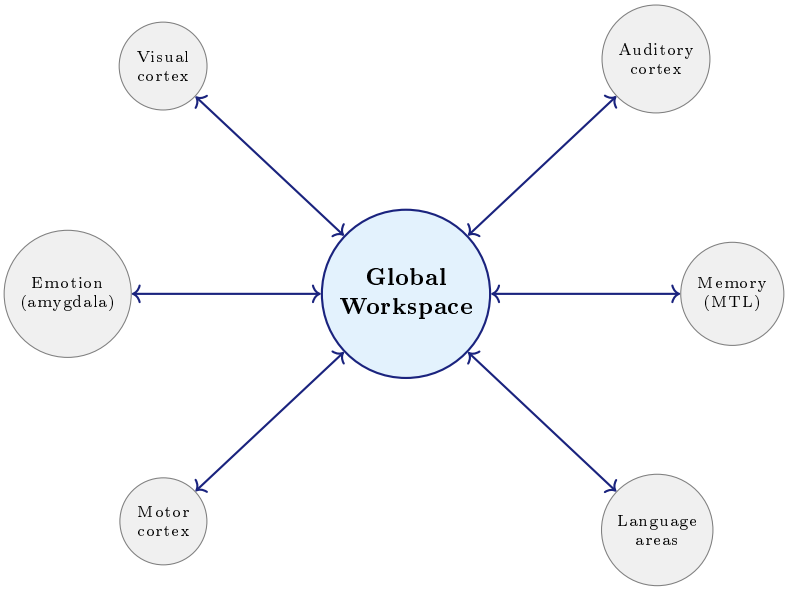

Figure: Hub-and-spoke architecture. The global workspace receives input from and broadcasts to all specialized processors.

4.2 The Ignition Dynamics: Mathematical Model

Dehaene and Changeux (2005) formalized ignition using a neural network model. The key is the interaction between local sensory processors and the global workspace, mediated by long-range excitatory connections.

4.2.1 Two-Population Model

Consider two coupled populations: a sensory population (S) and a workspace population (W). Their mean firing rates evolve according to:

$$\tau_S \frac{d\nu_S}{dt} = -\nu_S + \Phi_S(w_{SS}\,\nu_S + w_{SW}\,\nu_W + I_{\text{stim}}) \tag{4.1}$$

$$\tau_W \frac{d\nu_W}{dt} = -\nu_W + \Phi_W(w_{WW}\,\nu_W + w_{WS}\,\nu_S + I_{\text{top-down}}) \tag{4.2}$$

where:

- ν_S, ν_W: mean firing rates of sensory and workspace populations
- τ_S, τ_W: time constants (workspace often slower, $\tau_W > \tau_S$)
- w_{SS}: local recurrence in sensory area
- w_{WW}: recurrence within workspace (critical for ignition)
- w_{SW}, w_{WS}: long-range connections (bottom-up and top-down)
- $\Phi(\cdot)$: sigmoidal transfer function
- I_{stim}: sensory input; $I_{\text{top-down}}$: attentional modulation

4.2.2 Ignition as a Bifurcation

The key mathematical insight is that ignition corresponds to a **saddle-node bifurcation**. As I_{stim} increases:

1. For weak stimuli ($I_{\text{stim}} < I_{\text{crit}}$): the system has only one stable fixed point with $\nu_W \approx 0$ (no workspace activity). The stimulus is processed locally but **not** consciously perceived.
2. At the critical threshold ($I_{\text{stim}} = I_{\text{crit}}$): a saddle-node bifurcation occurs, creating a new stable fixed point with $\nu_W \gg 0$ (high workspace activity).
3. For strong stimuli ($I_{\text{stim}} > I_{\text{crit}}$): the system rapidly transitions to the high-activity state—**ignition**. The stimulus enters the workspace and is consciously perceived.

The critical threshold can be estimated from the fixed-point equation. At the bifurcation:

$$\frac{\partial}{\partial \nu_W}\left[-\nu_W + \Phi_W(w_{WW}\,\nu_W + \ldots)\right] = 0$$

which gives $w_{WW} \cdot \Phi'_W(\cdot) = 1$. This shows that ignition requires sufficiently strong recurrent workspace connectivity (w_{WW}) and/or sufficiently strong bottom-up input.

The all-or-none character of conscious perception emerges naturally: below threshold, $\nu_W \approx 0$ (unconscious); above threshold, ν_W jumps to a high value (conscious). There is no "partial" ignition—this is a nonlinear phase transition.

4.2.3 Top-Down Modulation and Attention

Top-down attention modulates the ignition threshold. In the model, $I_{\text{top-down}}$ effectively lowers I_{crit}, making it easier for stimuli to reach consciousness:

$$I_{\text{crit}}^{\text{attended}} < I_{\text{crit}}^{\text{unattended}}$$

This explains why attended stimuli are more easily perceived: attention pre-activates workspace neurons, bringing them closer to the ignition threshold.

4.3 Multi-Area Spiking Model

Dehaene, Changeux, and colleagues (2003) developed a more detailed model using spiking neurons organized into multiple cortical columns. Each column contains:

- 80% excitatory pyramidal neurons (further divided into layers)
- 20% inhibitory interneurons
- AMPA and NMDA excitatory synapses
- GABA_A inhibitory synapses

Columns are connected locally (within a cortical area) and long-range (via workspace neurons). The model implements:

$$C_m \frac{dV_i}{dt} = -g_L(V_i - V_L) - \sum_{\text{syn}} I_{\text{syn},i} + I_{\text{ext},i}$$

where V_i is the membrane potential of neuron i, g_L is the leak conductance, and I_{syn} includes AMPA, NMDA, and GABA$_A$ synaptic currents. Key findings from these simulations:

Simulation Results

1. **Subliminal processing:** Weak stimuli activate sensory columns transiently but fail to ignite workspace neurons. Activity dies out quickly.
2. **Ignition:** Sufficiently strong stimuli trigger a sudden, self-sustaining activation of workspace neurons across multiple areas, lasting several hundred milliseconds.
3. **Masking:** A second stimulus presented shortly after the first can prevent ignition by disrupting the bottom-up signal before it reaches the workspace—reproducing masking effects in psychophysics.
4. **Attentional blink:** While the workspace processes one stimulus, it is temporarily unavailable for a second stimulus, reproducing the attentional blink phenomenon.

Chapter 5

Formal Computational Models

5.1 The Broadcast Formalism

Let us formalize the broadcast operation. Consider a set of N modules $\{M_1, M_2, \ldots, M_N\}$, each with an internal state $\mathbf{x}_i(t) \in \mathbb{R}^{d_i}$. The workspace state is $\mathbf{w}(t) \in \mathbb{R}^{d_w}$.

Broadcast Dynamics

The dynamics of the system are:

$$\tau_i \frac{d\mathbf{x}_i}{dt} = F_i(\mathbf{x}_i) + G_i(\mathbf{w}) + \mathbf{u}_i(t) \tag{5.1}$$

$$\tau_w \frac{d\mathbf{w}}{dt} = F_w(\mathbf{w}) + \sum_{i=1}^{N} \alpha_i(t)\, H_i(\mathbf{x}_i) \tag{5.2}$$

where:

- F_i: internal dynamics of module i; F_w: internal workspace dynamics
- $G_i(\mathbf{w})$: broadcast signal from workspace to module i (top-down)
- $H_i(\mathbf{x}_i)$: input from module i to workspace (bottom-up)
- $\alpha_i(t) \in [0, 1]$: gating variable (attention/competition result)
- $\mathbf{u}_i(t)$: external input to module i

The **broadcast** is realized by the $G_i(\mathbf{w})$ term: once the workspace is activated (ignited), its state $\mathbf{w}$ is transmitted to *all* modules M_i, modulating their processing.

5.1.1 Competition via Gating

The gating variables $\alpha_i(t)$ implement competition. In the simplest case:

$$\alpha_i = \frac{e^{\beta s_i}}{\sum_j e^{\beta s_j}}$$

where $s_i = \|\mathbf{x}_i\| + b_i$ combines bottom-up signal strength $\|\mathbf{x}_i\|$ with top-down bias b_i, and β is an inverse temperature parameter. As $\beta \to \infty$, this becomes a hard winner-take-all: only the module with the highest s_i gains access ($\alpha_i \to 1$, all others $\to 0$).

5.2 Information-Theoretic Characterization

GWT's broadcast can be characterized information-theoretically.

5.2.1 Broadcast as Information Gain

Define the **information gain** from broadcast for module i as:

$$\Delta I_i = I(\mathbf{x}_i;\ S)\big|_{\text{after broadcast}} - I(\mathbf{x}_i;\ S)\big|_{\text{before broadcast}}$$

where S is the stimulus that entered the workspace. Before broadcast, module i has no information about S (if it was processed by a different module). After broadcast, it receives information via the workspace. The total information gain across all modules:

$$\Delta I_{\text{total}} = \sum_{i=1}^{N} \Delta I_i$$

This quantifies the "amplification" effect of consciousness: a single piece of information, once broadcast, becomes available across the entire brain.

5.2.2 Workspace Capacity and Information Bottleneck

The capacity limitation of the workspace can be formalized using the **information bottleneck** framework (Tishby et al., 1999). The workspace W compresses input information X while preserving information about the relevant variable Y:

$$\min_{p(w|x)} \Big[I(X;W) - \beta\, I(W;Y) \Big]$$

where $I(X;W)$ is the compression cost (limited by workspace capacity) and $I(W;Y)$ is the preserved relevant information. The parameter β trades off compression against information preservation. This formalization explains why consciousness is both limited (compressed) and informative (preserving relevant information).

5.3 Bayesian Formulation

GWT has also been connected to Bayesian inference. The workspace maintains a probability distribution over hypotheses, updated by Bayes' rule when new information enters:

$$p(H \mid D) = \frac{p(D \mid H)\, p(H)}{p(D)}$$

Conscious perception corresponds to the posterior $p(H \mid D)$ after sufficient evidence accumulation. The ignition threshold corresponds to a decision boundary in the Bayesian framework—the point at which the posterior probability of one hypothesis exceeds a criterion:

$$\text{Conscious access:} \quad p(H_1 \mid D) > \theta_{\text{crit}}$$

This connects GWT to accumulator models (drift-diffusion models) used in decision neuroscience. The decision variable $x(t)$ accumulates evidence:

$$dx = \mu\, dt + \sigma\, dW_t$$

where μ is the drift rate (signal strength) and $\sigma\, dW_t$ is noise. Conscious perception occurs when $x(t)$ first crosses the boundary θ. The mean reaction time is $\bar{T} = \theta/\mu$.

Chapter 6

Neural Signatures of Conscious Access

GNWT makes specific predictions about the neural signatures that distinguish conscious from unconscious processing. These have been extensively tested.

6.1 The P3b / P300 Wave

The most robust electrophysiological marker of conscious access is the **P3b component**—a positive-going ERP wave peaking ∼300–500 ms after stimulus onset at parietal electrodes.

P3b as Ignition Marker

According to GNWT:

- The P3b reflects the **ignition** of the global workspace.
- It is present for consciously perceived stimuli and **absent** for unperceived stimuli.
- Its latency (∼300 ms) reflects the time needed for bottom-up signals to reach and activate the workspace.
- Its widespread scalp topography reflects the global distribution of workspace neurons.
- Its amplitude correlates with subjective confidence in perception.

6.2 Late Sustained Activity

Beyond the P3b, conscious perception is associated with **late sustained activity** (>300 ms) in prefrontal and parietal areas, as measured by fMRI, MEG, and intracranial recordings. This sustained activity reflects the self-maintaining dynamics of the ignited workspace and the ongoing broadcast.

6.3 Gamma-Band Synchrony and Long-Range Connectivity

Conscious access is associated with increased **gamma-band** (30–100 Hz) oscillatory synchrony between distant brain areas:

Phase-locking value (PLV) between signals $x(t)$ and $y(t)$:

$$\mathrm{PLV} = \left| \frac{1}{T} \sum_{t=1}^{T} e^{i\,[\phi_x(t) - \phi_y(t)]} \right|$$

where $\phi_x(t)$ and $\phi_y(t)$ are the instantaneous phases extracted via the Hilbert transform or wavelet analysis. PLV ≈ 1 indicates perfect phase synchrony; PLV ≈ 0 indicates no consistent phase relationship.

GNWT predicts that conscious perception should show increased long-range PLV (especially in

beta/gamma bands), reflecting the coordinated activity of the ignited workspace. This has been confirmed experimentally.

6.4 The "All-or-None" Pattern

A distinctive prediction of GNWT is that the transition from unconscious to conscious processing is **nonlinear** (all-or-none), not gradual. This has been tested using:

1. **Masking paradigms:** Varying stimulus-mask SOA reveals a sharp threshold—below a critical SOA, perception is at chance; above it, perception is nearly perfect. Intermediate performance reflects trial-by-trial variability (ignition or no ignition), not graded awareness.
2. **fMRI studies:** Activity in prefrontal/parietal areas shows a bimodal distribution (high or low), not a continuum, correlated with reports of seeing vs. not seeing.
3. **Single-trial analysis:** Trial-by-trial neural signatures cluster into two distinct states (ignited/not ignited), with the probability of ignition varying smoothly with stimulus strength.

Chapter 7

Key Experimental Paradigms

7.1 Visual Masking

In **backward masking**, a brief target stimulus is followed by a mask after a variable stimulus-onset asynchrony (SOA). GWT predicts:

- At short SOAs (<50 ms): the mask disrupts bottom-up signals before ignition—target is invisible.
- At long SOAs (>100 ms): ignition has time to occur—target is seen.
- At intermediate SOAs: trial-by-trial variability (all-or-none ignition on each trial).

Dehaene et al. (2001) showed exactly this pattern using fMRI and ERPs, demonstrating bimodal brain responses consistent with ignition dynamics.

7.2 The Attentional Blink

When two targets (T1 and T2) are embedded in a rapid stream, T2 is often missed if it appears 200–500 ms after T1. GWT explains this as a **workspace refractory period**: while the workspace processes T1 (ignition + broadcast), it is temporarily unavailable for T2.

The mathematical model predicts the characteristic U-shaped function of T2 detection vs. T1–T2 lag:

- Lag 1 (very short): T2 can "sneak in" during T1 ignition (Lag-1 sparing)
- Lag 2–5 (∼200–500 ms): workspace occupied by T1, T2 is missed (blink)
- Lag 6+ (>500 ms): T1 processing complete, workspace available again

7.3 Binocular Rivalry

When different images are presented to each eye, perception alternates between them. GWT explains this as **alternating ignition**: each image competes for workspace access, and the currently conscious image maintains ignition until it is displaced by the competing image.

The alternation dynamics can be modeled with coupled oscillators:

$$\tau \frac{dx_1}{dt} = -x_1 + f(I_1 - \beta\, x_2 - a(t)\, x_1) \tag{7.1}$$

$$\tau \frac{dx_2}{dt} = -x_2 + f(I_2 - \beta\, x_1 - a(t)\, x_2) \tag{7.2}$$

where x_1, x_2 are activities of the two percepts, β is mutual inhibition, and $a(t)$ is an adaptation term that weakens the currently dominant percept over time, allowing the other to win.

The mean dominance duration $\bar{T}$ follows a gamma distribution, consistent with a noise-driven switching process:

$$p(T) = \frac{T^{k-1} e^{-T/\theta}}{\theta^k \, \Gamma(k)}$$

where k and θ are the shape and scale parameters.

7.4 Inattentional Blindness and Change Blindness

These phenomena demonstrate that even physically salient stimuli can fail to reach consciousness if attention is directed elsewhere—directly supporting the role of top-down attention in gating workspace access.

Chapter 8

Consciousness Under Altered States

8.1 General Anesthesia

GWT provides a natural explanation for the loss of consciousness under anesthesia. Anesthetic agents (e.g., propofol, sevoflurane) disrupt consciousness by:

1. **Disrupting long-range connectivity:** Anesthetics preferentially disrupt the long-range cortico-cortical connections that constitute the workspace, while leaving local processing relatively intact.
2. **Preventing ignition:** By reducing recurrent excitation (w_{WW}) and increasing inhibition, anesthetics raise the ignition threshold I_{crit} above any achievable sensory input.
3. **Breaking the broadcast:** Even if local sensory processing persists, the workspace cannot amplify and distribute the information.

Empirically, this is supported by studies showing that propofol specifically disrupts frontoparietal connectivity and prefrontal cortex feedback activity, while sensory-evoked responses in primary sensory cortex remain largely intact.

8.2 Sleep and Dreaming

NREM sleep: During deep sleep, cortical connectivity breaks down—neurons enter bistable "up" and "down" states, preventing sustained workspace ignition. GWT predicts reduced consciousness during NREM, consistent with the phenomenology.

REM sleep / Dreaming: During REM sleep, cortical connectivity is partially restored, and workspace ignition can occur—but driven by internal signals rather than external stimuli. This produces the vivid but often bizarre conscious experiences of dreaming.

8.3 Disorders of Consciousness

The GWT framework helps classify disorders of consciousness:

- **Vegetative state:** Intact local processing but no global ignition (workspace structurally or functionally disconnected).
- **Minimally conscious state:** Intermittent workspace ignition (partial recovery of long-range connectivity).
- **Locked-in syndrome:** Intact workspace and consciousness but disrupted motor output—the workspace is ignited, but the broadcast to motor systems is blocked.

Chapter 9

GWT vs. Other Theories

9.1 GWT vs. Integrated Information Theory (IIT)

Key Differences: GWT vs. IIT

Dimension	GWT / GNWT	IIT
What is consciousness?	Global broadcast/access	Integrated information (Φ)
Starting point	Function/mechanism	Phenomenology (axioms)
Hard Problem	Not directly addressed	Claimed to be addressed
Key brain area	Prefrontal-parietal	Posterior cortex
Key mechanism	Ignition + broadcast	Irreducible cause–effect power
Consciousness type	All-or-none access	Graded (Φ is continuous)
Panpsychism	No (requires complex architecture)	Yes ($\Phi > 0$ for simple systems)
Feedforward nets	Not conscious (no recurrent workspace)	Not conscious ($\Phi = 0$)
Testability	More empirically grounded	More mathematical/formal

The **Adversarial Collaboration** (Melloni et al., 2021) pits these theories against each other in pre-registered experiments. Key distinguishing predictions include:

- **Prefrontal involvement:** GWT predicts necessary prefrontal activation for consciousness; IIT predicts posterior cortex is sufficient.
- **All-or-none vs. graded:** GWT predicts all-or-none ignition; IIT allows graded consciousness.
- **Neural signatures:** GWT predicts late (>300 ms) prefrontal signatures; IIT predicts sustained posterior activity.

9.2 GWT vs. Higher-Order Theories

Higher-Order Theories (HOT, Rosenthal; HOROR, Lau) agree with GWT on the importance of prefrontal cortex but differ on mechanism: HOT requires a higher-order *representation about* the first-order state, while GWT requires broadcast/access. The two may converge: the workspace

broadcast might *be* the higher-order representation.

9.3 GWT vs. Recurrent Processing Theory

Recurrent Processing Theory (Lamme) agrees with GWT that recurrent processing is essential but disagrees on the locus: Lamme argues that local recurrent processing in sensory cortex suffices for (at least a form of) consciousness, without requiring prefrontal involvement. GWT insists that without global workspace ignition, there is no consciousness.

9.4 GWT vs. Predictive Processing

Predictive Processing (PP) and GWT are potentially complementary: PP describes the computational principles (prediction, prediction error) and GWT describes the architecture (global workspace) that implements them. The workspace might serve as the medium for global prediction-error propagation.

Chapter 10

Criticisms and Open Questions

10.1 Common Critiques

Critique 1: GWT does not address the Hard Problem.

GWT explains the *function* and *mechanisms* of consciousness but does not explain why global broadcast gives rise to subjective experience. Why isn't the workspace just a sophisticated information router with no "inner light"?

Response: GWT proponents argue that explaining the mechanisms *is* sufficient—the Hard Problem may dissolve once we fully understand the mechanisms. Alternatively, GWT is viewed as addressing the "real problem" (Seth, 2021): explaining the structure and contents of experience.

Critique 2: Overemphasis on prefrontal cortex.

Some evidence suggests that prefrontal activation during conscious perception may reflect *post-perceptual* processing (reporting, monitoring) rather than consciousness itself. The "no-report" paradigm (Frässle et al., 2014) shows that when report-related confounds are removed, prefrontal activation is greatly reduced.

Response: GNWT proponents argue that no-report paradigms don't fully eliminate workspace involvement—even without verbal report, the workspace may ignite for internal monitoring. This remains actively debated.

Critique 3: The consciousness/access distinction.

Ned Block distinguishes between **phenomenal consciousness** (subjective experience) and **access consciousness** (availability for report and reasoning). Block argues GWT only explains access consciousness, leaving phenomenal consciousness unaddressed. A rich perceptual experience may be "overflowing" the workspace capacity.

Response: Dehaene and colleagues argue that phenomenal consciousness without access is incoherent—if information is not accessed, there is no evidence it is conscious.

Critique 4: Underspecification of workspace boundaries.

The theory does not precisely define which neurons are "workspace neurons" vs. "peripheral processors." The boundary is fuzzy.

Response: The workspace is defined functionally (neurons with long-range connections that participate in ignition/broadcast), not anatomically. The boundary is inherently graded, not sharp.

10.2 Open Questions

1. **What is the exact capacity of the workspace?** Is it one item, a few chunks, or a more complex structure?

2. **How does the workspace interact with working memory?** Are they the same system or different?
3. **Can multiple ignitions coexist?** Or is the workspace strictly serial?
4. **What role does thalamocortical interaction play?** The thalamus may be essential for workspace ignition but is underspecified in most models.
5. **Can GWT be reconciled with IIT?** Some theorists (e.g., Mashour, Alkire) have proposed hybrid frameworks combining workspace ignition with integrated information.
6. **How does GWT apply to non-human animals?** Do animals with less developed prefrontal cortex have a workspace? Do they have consciousness?

Chapter 11

Further Reading and References

Core GWT Publications

[1] Baars, B. J. (1988). *A Cognitive Theory of Consciousness*. Cambridge University Press.

[2] Baars, B. J. (1997). *In the Theater of Consciousness: The Workspace of the Mind*. Oxford University Press.

[3] Baars, B. J. (2005). Global workspace theory of consciousness: toward a cognitive neuroscience of human experience. *Progress in Brain Research*, 150, 45–53.

[4] Dehaene, S., Kerszberg, M., & Changeux, J.-P. (1998). A neuronal model of a global workspace in effortful cognitive tasks. *Proceedings of the National Academy of Sciences*, 95(24), 14529–14534.

[5] Dehaene, S. & Changeux, J.-P. (2011). Experimental and theoretical approaches to conscious processing. *Neuron*, 70(2), 200–227.

Neural Mechanisms and Ignition

[6] Dehaene, S., Sergent, C., & Changeux, J.-P. (2003). A neuronal network model linking subjective reports and objective physiological data during conscious perception. *PNAS*, 100(14), 8520–8525.

[7] Dehaene, S., et al. (2006). Conscious, preconscious, and subliminal processing: a testable taxonomy. *Trends in Cognitive Sciences*, 10(5), 204–211.

[8] Sergent, C., Baillet, S., & Dehaene, S. (2005). Timing of the brain events underlying access to consciousness during the attentional blink. *Nature Neuroscience*, 8(10), 1391–1400.

[9] van Vugt, B., et al. (2018). The threshold for conscious report: Signal loss and response bias in visual and frontal cortex. *Science*, 360(6388), 537–542.

Experimental Evidence

[10] Dehaene, S., et al. (2001). Cerebral mechanisms of word masking and unconscious repetition priming. *Nature Neuroscience*, 4(7), 752–758.

[11] Del Cul, A., Baillet, S., & Dehaene, S. (2007). Brain dynamics underlying the nonlinear threshold for access to consciousness. *PLoS Biology*, 5(10), e260.

[12] Mashour, G. A., Roelfsema, P., Changeux, J.-P., & Dehaene, S. (2020). Conscious processing and the global neuronal workspace hypothesis. *Neuron*, 105(5), 776–798.

Adversarial Collaboration and Theory Comparison

[13] Melloni, L., et al. (2021). Making the hard problem of consciousness easier. *Science*, 372(6545), 911–912.

[14] Cogitate Consortium (2023). An adversarial collaboration to critically evaluate theories of consciousness. Preregistered report.

[15] Block, N. (2005). Two neural correlates of consciousness. *Trends in Cognitive Sciences*, 9(2), 46–52.

[16] Seth, A. K. (2021). *Being You: A New Science of Consciousness*. Dutton.

Computational Models and Mathematics

[17] Wilson, H. R. & Cowan, J. D. (1972). Excitatory and inhibitory interactions in localized populations of model neurons. *Biophysical Journal*, 12(1), 1–24.

[18] Tishby, N., Pereira, F., & Biller, W. (1999). The information bottleneck method. *Proc. 37th Allerton Conference*.

[19] Deco, G., et al. (2015). Rethinking segregation and integration: contributions of whole-brain modelling. *Nature Reviews Neuroscience*, 16, 430–439.

[20] Shanahan, M. (2010). *Embodiment and the Inner Life: Cognition and Consciousness in the Space of Possible Minds*. Oxford University Press.

Related Topics

[21] Fodor, J. A. (1983). *The Modularity of Mind*. MIT Press.

[22] Baddeley, A. (2000). The episodic buffer: a new component of working memory? *Trends in Cognitive Sciences*, 4(11), 417–423.

[23] Lamme, V. A. F. (2006). Towards a true neural stance on consciousness. *Trends in Cognitive Sciences*, 10(11), 494–501.

[24] Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated information theory: an updated account. *Archives Italiennes de Biologie*, 154, 1–18.

End of Document